

Appendix of “Automated Random Embedding for Practical Bayesian Optimization with Unknown Effective Dimension”

Hong Qian^{1*}, Xiang Shu², Xiang Xia¹, Xuhui Liu³, Yangde Fu¹, Bei Liang¹,
Huibin Wang¹ and Liang Dou¹

¹Shanghai Institute of AI for Education, and School of Computer Science and Technology,
East China Normal University, Shanghai 200062, China

²Ant Group, Hangzhou 310013, China


³Nanjing University, Nanjing 210023, China

hqian@cs.ecnu.edu.cn, shuxiang.shu@antgroup.com, xxia@stu.ecnu.edu.cn,
liuxh@lamda.nju.edu.cn, {ydfu, bliang, bingowang}@stu.ecnu.edu.cn, ldou@cs.ecnu.edu.cn

This appendix of “Automated Random Embedding for Practical Bayesian Optimization with Unknown Effective Dimension” is organized as follows. We begin with detailed preliminaries, followed by the pseudo-code of DSEBO and the dataset initialization algorithm. We then provide proofs of the theoretical results in the main paper. Next, we describe the real-world datasets used in our experiments and the implementation details of the comparative methods. Finally, we present detailed experimental results along with a comprehensive hyper-parameter analysis.

A Detailed Preliminaries

This section provides additional details on Bayesian optimization, the optimal ϵ -effective dimension, and random embedding to further clarify the preliminaries, necessary background and notation.



(a) Function with effective dimension (b) ϵ -effective dimension for MSLR

Figure 1: An illustration for a synthetic function with effective dimension and the test for the existence of ϵ -effective dimension in real dataset MSLR Zhang *et al.* [2023].

A.1 Bayesian Optimization

Bayesian optimization (BO) [Srinivas *et al.*, 2010; Shahriari *et al.*, 2016; Garnett, 2023] is a well-known derivative-free optimization method that strategically utilizes prior knowledge to guide the sampling process. First, BO constructs a surrogate model, typically a Gaussian process (GP), based on

samples of the objective function. Next, BO calculates the acquisition function based on this model to guide the subsequent sampling process and the trade-off between “exploration” and “exploitation” in the search space.

Based on the observed dataset $\mathcal{D}$, BO estimates the posterior distribution of the objective function $P(f|\mathcal{D}) \propto P(\mathcal{D}|f)P(f)$, where $P(f)$ denotes the prior distribution of the objective function and $P(\mathcal{D}|f)$ is the likelihood. This posterior distribution incorporates information about the objective function and is used to inform subsequent modeling and optimization. Once the posterior distribution has been obtained, BO employs an acquisition function to determine the next sampling point. A common choice is UCB [Srinivas *et al.*, 2010], defined as $\alpha_{\text{UCB}}(\mathbf{x}) = \mu_t(\mathbf{x}) + \sqrt{\kappa_{t+1}}\sigma_t(\mathbf{x})$, where $\mu_t(\mathbf{x})$ estimates the objective function, associated with “exploitation”, while $\sigma_t(\mathbf{x})$ represents the uncertainty of the objective function, associated with “exploration”. The hyper-parameter κ in UCB balances “exploration” and “exploitation” for efficient optimization.

A.2 Random Embedding

If the high-dimensional objective function has effective dimension, the subspace embedding technique can effectively achieve dimensionality reduction. For example, Figure 1(a) shows a 2-dimensional synthetic function $f(x_1, x_2)$ with a 1-dimensional effective dimension. In such cases, only a few dimensions significantly influence the objective function and are prioritized during optimization.

However, it is challenging to satisfy the effective dimension assumption in real tasks. In this case, Qian *et al.* [2016] propose the concept of optimal ϵ -effective dimension to relax the assumption.

Empirical evidence shows that many high-dimensional tasks [Sun *et al.*, 2022] and datasets [Zhang *et al.*, 2023] follow the ϵ -effective dimension. A typical example is the MSLR task (as Figure 1(b)), where 50-dimensional subspace fully captures the variation in the objective function, and other dimensions have a low impact, i.e., the ϵ -effective dimension of the MSLR dataset does not exceed 50.

Random embedding (RE) [Wang *et al.*, 2016; Nayebi *et al.*, 2019] is a widely used subspace embedding method as shown in Figure 2. For a high-dimensional optimization problem

*Corresponding Author.




Figure 2: The process of BO with random embedding.

$\mathbf{x}^* = \operatorname{argmin}_{\mathbf{x} \in \mathcal{X} \subset \mathbb{R}^D} f(\mathbf{x})$, a random matrix $A \in \mathbb{R}^{D \times d}$ with elements sampled from a Gaussian distribution $\mathcal{N}(0, \sigma^2)$ embeds the d -dimensional subspace $\mathcal{Z} \subset \mathbb{R}^d$ into the D -dimensional search space $\mathcal{X} \subset \mathbb{R}^D$. The optimizer only needs to optimize $\mathbf{z} \in \mathcal{Z}$ in the low-dimensional subspace, and then embed $\mathbf{z}$ into $\mathcal{X}$ through the matrix A to get the solution $\mathbf{x} = p_{\mathcal{X}}(A\mathbf{z})$. Here, $p_{\mathcal{X}}(A\mathbf{z}) = \operatorname{argmin}_{\mathbf{x} \in \mathcal{X}} \|\mathbf{x} - A\mathbf{z}\|_2$ is to project illegal points (i.e., $A\mathbf{z} \notin \mathcal{X}$) back to the $\mathcal{X}$ region. Subsequently, the objective function f is evaluated in the solution $\mathbf{x}$, and the sample point $(\mathbf{z}, f(p_{\mathcal{X}}(A\mathbf{z})))$ is added to the subspace dataset. The optimizer uses the updated dataset to complete the next optimization iteration.

B Pseudo-code of DSEBO

The pseudo-code of DSEBO is shown in Algorithm 1. The algorithm aims to optimize a high-dimensional objective function $f(\cdot)$ with unknown effective dimensions on a D -dimensional space within a limited number of evaluations to find its minimum value. DSEBO requires *budget* (i.e., total evaluations), the subspace dimension search range $[d_l, d_h]$, and a hyper-parameter β . For a high-dimensional problem, one can search from the dimension 5, i.e., $d_l = 5$, and the recommended value of the upper boundary of the search range d_h is $\min(D, 100)$. When the dimension exceeds 100, the performance of the Bayesian optimizer will drop sharply, so the search range of the subspace will not exceed 100. The hyper-parameter β controls the scale of the dimension expansion.

Before optimization, the algorithm will initialize the shared embedding matrix and set some necessary parameters, as shown in lines 1–2, where T controls the judgment of whether the solution converges, indirectly controlling the number of dimension expansion iterations. During the optimization phase, if a subspace has not been initialized, the algorithm will initialize its dataset, optimizer, current optimal solution b , and random embedding matrix in sequence, as shown in lines 4–9. Then, the following step is Bayesian optimization based on stochastic embedding, including steps such as the optimizer giving the next solution, embedding the solution into the high-dimensional search space, querying the objective function, updating the dataset and the optimizer, as shown in lines 10–14. The updating of the optimal value b achievable by the current algorithm occurs in lines 15–17, with the error constrained within the range of α . If the solution has converged in the current subspace, i.e., b has not changed in T iterations, expand to the new subspace and update T (the convergence

time T of the higher-dimensional subspace solution is usually larger, so we update it), as shown in lines 18–20. Repeat the above steps until the *budget* is exhausted, and finally embed the subspace into the search space to obtain the optimal value found by the algorithm.

C Pseudo-code of Dataset Initialization Algorithm

In high-dimensional optimization problems, when the optimization process enters a new subspace, it is necessary to initialize the dataset in the new dimensional space. To fully utilize the data from the existing lower-dimensional subspace, this algorithm generates the new dataset by extending the existing data into the higher-dimensional space. As shown in Algorithm 2, in the first iteration ($t = 1$), the dataset starts empty and is initialized with a random point in the current dimensional space. For subsequent dimension expansions, the dataset $\mathcal{D}^{(t)}$ for the current dimension is generated by expanding the previous dataset $\mathcal{D}^{(t-1)}$. This expansion is achieved by appending zeros to each solution vector $\mathbf{z}$, thereby extending it into the new subspace while preserving the original evaluation

Algorithm 1 DSEBO Algorithm

Input: D -dimensional objective function $f(\cdot)$, *budget*, $d_l = 5$, $d_h = \min(D, 100)$, hyper-parameter β .

- 1: Initialize the shared embedding matrix $S_{D \times d_h}$, whose elements are sampled from $\mathcal{N}(0, \sqrt{1/d_h})$
- 2: $t = 1, d^{(1)} = d_l, T = \lfloor \frac{\text{budget}}{2\beta} \rfloor$
- 3: **while** $t \leq \text{budget}$ **do**
- 4: **if** $t = 1$ or $d^{(t)} \neq d^{(t-1)}$ **then**
- 5: $\mathcal{D}^{(t)} \leftarrow$ Initialize dataset using shared embedding
- 6: Using $\mathcal{D}^{(t)}$ update the optimizer $B_{d^{(t)}}$
- 7: $b \leftarrow \min\{y \mid (z, y) \in \mathcal{D}^{(t)}\}$
- 8: $A_{d^{(t)}} \leftarrow$ The first $d^{(t)}$ rows of matrix $S_{D \times d_h}$
- 9: **end if**
- 10: $\mathbf{z}^{(t)} \leftarrow$ Next solution given by the optimizer $B_{d^{(t)}}$
- 11: $\mathbf{x}^{(t)} \leftarrow$ Random Embedding $\mathbf{z}^{(t)}$, i.e., $p_{\mathcal{X}}(A_{d^{(t)}} \mathbf{z}^{(t)})$
- 12: $f(\mathbf{x}^{(t)}) \leftarrow$ Query the objective function
- 13: $\mathcal{D}^{(t+1)} \leftarrow \mathcal{D}^{(t)} \cup \{(\mathbf{z}^{(t)}, f(\mathbf{x}^{(t)}))\}$
- 14: Using $\mathcal{D}^{(t+1)}$ update the optimizer $B_{d^{(t+1)}}$
- 15: **if** $|f(\mathbf{x}^{(t)}) - b| > 0.5$ **then**
- 16: $b \leftarrow f(\mathbf{x}^{(t)})$
- 17: **end if**
- 18: **if** b has no change in T iterations **then**
- 19: $d^{(t+1)} \leftarrow$ Update dimension using dynamic dimension expanding strategy
- 20: $T \leftarrow \lfloor (1 + \frac{d^{(t)} - d_l}{d_h - d_l}) \frac{\text{budget}}{2\beta} \rfloor$
- 21: **else**
- 22: $d^{(t+1)} \leftarrow d^{(t)}$
- 23: **end if**
- 24: $t \leftarrow t + 1$
- 25: **end while**
- 26: $\mathbf{z}^* \leftarrow \arg \max_{\mathbf{z}} \{y \mid (\mathbf{z}, y) \in \mathcal{D}^{(t)}\}$

Output: The best solution $\mathbf{x}^* = A^{(t)} \mathbf{z}^*$.

Algorithm 2 Dataset Initialization

Input: Previous dimension $d^{(t-1)}$, current dimension $d^{(t)}$, dataset $\mathcal{D}^{(t-1)}$ of previous space.

```
1: if  $t = 1$  then  
2:    $\mathcal{D}^{(1)} \leftarrow$  Add a random  $d^{(1)}$ -dimensional solution  
3: else  
4:    $\mathcal{D}^{(t)} \leftarrow \emptyset$   
5:   for all  $(z, y) \in \mathcal{D}^{(t-1)}$  do  
6:      $z' = (z^\top, \mathbf{0}_{d^{(t)}-d^{(t-1)}})^\top$   
7:      $\mathcal{D}^{(t)} \leftarrow \mathcal{D}^{(t)} \cup \{(z', y)\}$   
8:   end for  
9: end if
```

Output: Dataset $\mathcal{D}^{(t)}$ of current space.

value y . This method ensures that valuable data from lower dimensions are retained and utilized in the higher-dimensional optimization process.

D Proof of Section 5

D.1 Proof of Theorem 5.1

Proof. By Theorem 2 of Srinivas *et al.* [2010], the regret bound is $O^*\left(\sqrt{T(\log T)^{d+1}}\right)$ in the absence of approximation error. To complete the analysis, we need to account for the impact of the approximation error.

From Lemma 1 of Qian *et al.* [2016], the approximation error is given by $2\epsilon_d$. Combining these results, the total regret bound becomes the sum of the regret without approximation error and the approximation error term. This completes the proof. $\square$

D.2 Proof of Theorem 5.2

Proof. Let $\sigma^{(t)}(\cdot)$ be the standard deviation function in step t , $r^{(t)}$ be the single-step regret without considering the approximation error. Lemma 5.2 of Srinivas *et al.* [2010] reveals that

$(r^{(t)})^2 = 4\beta^{(t)}(\sigma^{(t-1)})^2(\mathbf{x}^{(t-1)})$. Therefore,

$$\begin{aligned} & \sum_{t=1}^T (r^{(t)})^2 \\ &= \sum_{t=1}^T 4\beta^{(t)}(\sigma^{(t-1)})^2 \\ &\stackrel{(a)}{\leq} \sum_{t=1}^T 4\beta^{(T)}(\sigma^{(t-1)})^2 \\ &= 4\beta^{(T)} \sum_{t=1}^T (\sigma^{(t-1)})^2 (\mathbf{x}^{(t-1)}) \\ &= 4\beta^{(T)} \sum_{t=1}^T \sigma^2 \left(\sigma^{-2}(\sigma^{(t-1)})^2 (\mathbf{x}^{(t-1)}) \right) \\ &\leq 4\beta^{(T)} \sum_{t=1}^T \sigma^2 C_2 \log \left(1 + \sigma^{-2}(\sigma^{(t-1)})^2 (\mathbf{x}^{(t-1)}) \right) \\ &\stackrel{(b)}{\leq} C_1 \beta^{(T)} \sum_{h=1}^H \gamma_{T_h}, \end{aligned}$$

where (a) is because $\beta^{(t)}$ is nondecreasing, (b) comes from Lemma 5.3 of Srinivas *et al.* [2010].

According to Theorem 5 of Srinivas *et al.* [2010], we have $\gamma_{T_h} = O(\log T_h)^{d_h+1}$. Plugging this into the above inequality, and use Cauchy-Schwarz inequality as Lemma 5.4 in Srinivas *et al.* [2010], we obtain the final result. $\square$

E Real-World Datasets

In this section, we provide a detailed introduction to the real-world datasets used in our experiments.

MSLR [Qin *et al.*, 2010]: The first real-world dataset is the Microsoft Learning to Rank (MSLR) dataset [Qin *et al.*, 2010], specifically the MSLR-WEB-10K version, which includes over 10000 queries, each with 136 features per website page. Each website page has a relevance judgment, ranging from 0 (irrelevant) to 4 (completely relevant). When dealing the dataset, we normalize the dataset and then use neural networks to model the dataset as the objective function. Given the complexity of using the dataset as an oracle, we assess preferences through the neural network’s predictions instead. The architecture of this neural network is designed with three hidden layers, each containing 128, 64, and 32 neurons, and utilizes the Sigmoid function as the activation mechanism.

Lasso-Hard [Sehic *et al.*, 2022]: The second dataset is Lasso-Hard from LassoBench [Sehic *et al.*, 2022], a 1000-dimensional optimization task designed to identify sparse regression coefficients that minimize Lasso regression loss.

LIMO [Eckmann *et al.*, 2022]: The third dataset is LIMO [Eckmann *et al.*, 2022], a framework for molecular generation aimed at optimizing specific properties by operating in a 1024-dimensional latent space learned through a variational autoencoder. Specifically, the LIMO dataset is a drug-like molecule design task, where the objective function is designed to balance the octanol-water partition coefficient, accessibility, and the presence of large rings.




Figure 3: Results on remaining 1000-dimensional synthetic functions compared with various high-dimensional optimization algorithms. All algorithms are independently repeated 10 times.

F Baselines Implementation Details

F.1 High-dimensional Optimization Methods

REMBO [Wang *et al.*, 2016]: REMBO, the first random embedding approach, repeats using the BoTorch framework in the experiments. We adhere to the same hyper-parameter specifications detailed in [Wang *et al.*, 2016].

HesBO [Nayebi *et al.*, 2019]: HesBO takes a novel approach to avoid embedding at the boundary by altering the generation process of the embedding matrix. For our implementation, we have utilized the version made available by the author at the GitHub repository: <https://github.com/aminnayebi/HeSBO>.

SIRBO [Zhang *et al.*, 2019]: Unlike traditional random embedding techniques, SIRBO computes the embedding matrix using the sliced inverse regression (SIR) method. We use the author’s implementation from https://github.com/cjfcstj/SILBO/blob/master/SIR_BO.py.

VAEBO [Gómez-Bombarelli *et al.*, 2018]: VAE-BO employs a variational auto-encoder (VAE) to discern the embedding relationship between high-dimensional spaces and their lower-dimensional counterparts. And we also utilize the code made available by the author: https://github.com/lamda-bbo/MCTS-VS/blob/master/baseline/vae_bo.py, adjusting the learning rate to 0.001 and updating the VAE model every 20 iterations.

ALEBO [Letham *et al.*, 2020]: ALEBO refines the acquisition function within constraints to become the state-of-the-art (SOTA) method for random embedding. We use the author’s implementation from <https://github.com/facebookresearch/alebo>.

BAXUS [Papenmeier *et al.*, 2022]: BAXUS is an embedding-based method designed to optimize high-dimensional black-box functions by using nested random subspaces and a unique dimensionality growth strategy. We use the author’s code from <https://github.com/LeoIV/BAXUS>.

MCTS-VS [Song *et al.*, 2022]: MCTS-VS employs Monte Carlo tree search to iteratively select and optimize a subset of variables within a low-dimensional subspace. The implementation from <https://github.com/lamda-bbo/MCTS-VS> is used.

TuRBO [Eriksson *et al.*, 2019]: TuRBO is an efficient method for handling high-dimensional optimization by dividing the space into smaller regions for local optimization. We use the code from <https://github.com/uber-research/TuRBO>.

SAASBO [Eriksson and Jankowiak, 2021]: SAASBO focuses on optimizing only a few important dimensions, thus reducing computational complexity in high-dimensional spaces. We implement the method using the author’s code from <https://github.com/martinjankowiak/saasbo>.

DuMBO Bardou *et al.* [2024]: DuMBO employs decentralized message-passing and a refined acquisition function to relax additive structure constraints in high-dimensional Bayesian Optimization, achieving asymptotic optimality on functions with complex decompositions. The code for DuMBO can be found at <https://github.com/abardou/dumbo>.

RDUCB [Ziomek and Bou-Ammar, 2023]: RDUCB uses random tree decompositions to construct additive GP models with cycle-free pairwise dimensional interactions, effectively addressing the challenge of being misled by local data in high-dimensional optimization. The code for RDUCB is available at <https://github.com/huawei-noah/HEBO>.

SBO-SE [Xu *et al.*, 2025]: SBO-SE employs a robust strategy to initialize the length-scale of the GP kernel, avoiding the vanishing gradient problem during Gaussian process training in high-dimensional Bayesian optimization. The code we use is implemented based on the BoTorch library.

LMMAES Loshchilov *et al.* [2019]: LMMAES reduces the time and space complexity of traditional matrix adaptation evolution strategy by approximating the covariance structure using a small set of evolution paths. The code for LMMAES is available as part of the pypop library at <https://github.com/Evolutionary-Intelligence/pypop>.

DCEM Amos and Yarats [2020]: DCEM is a differentiable variant of the cross-entropy method that enables gradient-based end-to-end learning by employing a smooth top-k operation, addressing the challenges of backpropagating through non-differentiable or discrete optimization steps. The implementation of DCEM can be found in the pypop library at <https://github.com/Evolutionary-Intelligence/pypop>.

F.2 Multi-armed Bandit Strategy

ϵ -greedy [Langford and Zhang, 2007]: This method strikes a simple balance between exploration and exploitation by selecting the best known action most of the time while occasionally choosing randomly with a small probability ϵ . This ensures that the algorithm does not rely solely on the existing knowledge and periodically explores other options, potentially discovering more optimal strategies. In our experiments, we set $\epsilon = 0.5$.

C-UCB [Auer *et al.*, 2002]: Classic Upper Confidence Bound (C-UCB) is the foundational algorithm in the UCB family, which selects actions based on a trade-off between their past rewards and a confidence interval, ensuring a balance between exploiting known rewards and exploring less certain options. In our experiments, the κ is set to $\sqrt{2 \log(T)/n_i}$ as default, where T is the total number of iterations and n_i is the number of i -th subspace optimized time.

UCB-E [Audibert *et al.*, 2010]: The UCB-E algorithm enhances the traditional Upper Confidence Bound approach by introducing a controllable exploration factor, c . This factor allows for fine-tuning the exploration level, independent of each arm’s estimated uncertainty, making it especially useful in




Figure 4: Results on 10000-dimensional synthetic functions compared with various high-dimensional optimization algorithms. All algorithms are independently repeated 10 times.




Figure 5: Results on synthetic functions compared with different MAB strategies, and the subspace dimensions selected by DSEBO throughout the optimization process. All algorithms are independently repeated 10 times.

scenarios where the standard uncertainty model might not adequately represent the exploratory needs. In our experiments, we set $c = 0.5$.

TS [Jin *et al.*, 2022]: Thompson Sampling (TS) is a probabilistic approach that selects arms based on samples drawn from their estimated reward distributions, effectively balancing

exploration and exploitation by adapting to the uncertainty in the estimate of each arm's reward.

Softmax [Sutton and Barto, 2018]: In the Softmax method, the selection of arms is controlled by a probability distribution weighted by the estimated value of each arm, which is adjusted according to the temperature parameter τ . Lower τ favors the




Figure 6: Results on the 1000-dimensional Sphere function with an effective dimension of 30 for varying c , compared with various algorithms. All algorithms are independently repeated 10 times.




Figure 7: Results on the 1000-dimensional synthetic function with varying effective dimension d_e . All algorithms are independently repeated 10 times.




Figure 8: Results of hyper-parameter β experiments in the synthetic functions with $D = 1000$, $d_e = 30$. All experiments are repeated 10 times.

best arms, while higher τ increases exploration of unknown arms. In our experiments, we set $\tau = 1.0$.

S-Halving [Karnin *et al.*, 2013]: Successive Halving is a method of saving budgets by first allocating equal resources to all candidates and then phasing out underperforming candidates, focusing resources on the most promising bandits as the process iterates.

Extreme: The extreme strategy is designed to achieve exploitation by aggressively focusing on the arms that appear to be optimal, often used in environments where the cost of exploration is high or when quick decisions are critical.

Expectation: The expectation strategy selects arms based on the expected value of the reward, typically favoring arms with higher average rewards.

Random: The random strategy selects actions randomly, ensuring equal exploration of all available options. It is useful for baseline comparisons or when no prior data exists. Although simple, it can occasionally reveal overlooked possibilities in complex scenarios.

G Detailed Results

The remaining experimental results of the synthetic functions with $D = 1000$ mentioned in the Section 6 are shown in Figure 3. The experimental results of the synthetic functions with $D = 10000$ are shown in Figure 4. When dealing with high-dimensional tasks with $D = 10000$, since SIRBO cannot complete 500 iterations within 8 hours, only the results of the first 100 iterations are plotted. Experimental results show that on 6 synthetic functions, DSEBO shows the fastest convergence speed and can find solutions with high performance.

Additionally, the comparison results of synthetic functions between DSEBO and a series of MAB strategies are presented in Figure 5. As evidenced by the figures, DSEBO verifies significant advantages over other arm selection strategies by dynamically expanding subspace dimensions, efficiently sharing data across dimensions, and achieving faster convergence and superior performance in high-dimensional optimization tasks.

Table 1, Table 2, Table 3, Table 4 and Table 5 record the final mean convergence value of various algorithms under each experimental environment, the optimal solution that can be




Figure 9: Results of hyper-parameter d_h experiments in the synthetic functions with $D = 1000$, $d_e = 30$. All experiments are repeated 10 times.

found, and the mean operation time. The results show that DSEBO can provide a better dynamic dimension expanding strategy and show good optimization performance in all synthetic functions and real-world tasks, which reflects in stable convergence performance and the ability to explore excellent solutions.

To further evaluate the capability of DSEBO, we also conduct experiments by adjusting the shifting scalar c and the effective dimension d_e of the synthetic functions, as shown in Figure 6 and Figure 7, which verifies that DSEBO consistently maintains its superior performance across different settings.

H Hyper-parameter Analysis

We conduct hyper-parameter analysis on β and the upper boundary of search range d_h on two synthetic functions to verify the robustness of DSEBO under different hyper-parameter configurations, as shown in Figure 8 and Figure 9.

The experiments of hyper-parameter β show that when β is too large (e.g., $\beta = 24, 32$), the dimension updates more frequently, but the change in dimensionality is very small. Conversely, when β is too small (e.g., $\beta = 1, 4$), the change in dimensionality becomes less intuitive, significantly affecting the convergence of the solution. Nevertheless, DSEBO consistently finds good solutions across different β configurations, except for the extreme hyper-parameter values (e.g., $\beta = 1, 4$). From the experiments of the upper boundary of search range d_h , it can be observed that setting $d_h = 100$ yields satisfactory optimization performance. Moreover, compared to $d_h = 80$ and $d_h = 120$, $d_h = 100$ results in a more reasonable subspace dimension selection that does not simply opt for higher dimensions indiscriminately.

The hyper-parameter analysis verifies that the chosen hyper-parameter values are reasonable, and shows the robustness of the DSEBO across different settings.

References

Brandon Amos and Denis Yarats. The differentiable cross-entropy method. In *Proceedings of the 37th International Conference on Machine Learning*, volume 119, pages 291–302, Virtual Event, 2020.

Jean-Yves Audibert, Sébastien Bubeck, and Rémi Munos. Best arm identification in multi-armed bandits. In *Proceedings of the 23rd Conference on Learning Theory*, pages 41–53, Haifa, Israel, 2010.

Peter Auer, Nicolò Cesa-Bianchi, and Paul Fischer. Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47:235–256, 2002.

Anthony Bardou, Patrick Thiran, and Thomas Begin. Relaxing the additivity constraints in decentralized no-regret high-dimensional Bayesian optimization. In *Proceedings of the 12th International Conference on Learning Representations*, 2024.

Peter Eckmann, Kunyang Sun, Bo Zhao, Mudong Feng, Michael K. Gilson, and Rose Yu. LIMO: latent inceptionism for targeted molecule generation. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162, pages 5777–5792, Baltimore, MD, 2022.

David Eriksson and Martin Jankowiak. High-dimensional Bayesian optimization with sparse axis-aligned subspaces. In *Proceedings of the 37th Conference on Uncertainty in Artificial Intelligence*, pages 493–503, Virtual Event, 2021.

David Eriksson, Michael Pearce, Jacob R. Gardner, Ryan Turner, and Matthias Poloczek. Scalable global optimization via local Bayesian optimization. In *Advances in Neural Information Processing Systems 32*, pages 5497–5508, Vancouver, Canada, 2019.

Roman Garnett. *Bayesian Optimization*. Cambridge University Press, Cambridge, 2023.

Rafael Gómez-Bombarelli, Jennifer N Wei, David Duvenaud, José Miguel Hernández-Lobato, Benjamín Sánchez-Lengeling, Dennis Sheberla, Jorge Aguilera-Iparraguirre, Timothy D Hirzel, Ryan P Adams, and Alán Aspuru-Guzik. Automatic chemical design using a data-driven continuous representation of molecules. *ACS Central Science*, 4(2):268–276, 2018.

Tianyuan Jin, Pan Xu, Xiaokui Xiao, and Anima Anandkumar. Finite-time regret of Thompson sampling algorithms for exponential family multi-armed bandits. In *Advances in Neural Information Processing Systems 35*, pages 38475–38487, New Orleans, LA, 2022.

Zohar Shay Karnin, Tomer Koren, and Oren Somekh. Almost optimal exploration in multi-armed bandits. In *Proceedings of the 30th International Conference on Machine Learning*, pages 1238–1246, Atlanta, GA, 2013.

John Langford and Tong Zhang. The epoch-greedy algorithm for multi-armed bandits with side information. In *Advances in Neural Information Processing Systems 20*, pages 817–824, Vancouver, Canada, 2007.

Benjamin Letham, Roberto Calandra, Akshara Rai, and Eytan Bakshy. Re-examining linear embeddings for high-dimensional Bayesian optimization. In *Advances in Neural Information Processing Systems 33*, Virtual Event, 2020.

Ilya Loshchilov, Tobias Glasmachers, and Hans-Georg Beyer. Large scale black-box optimization by limited-memory matrix adaptation. *IEEE Transactions on Evolutionary Computation*, 23(2):353–358, 2019.

Amin Nayebi, Alexander Munteanu, and Matthias Poloczek. A framework for Bayesian optimization in embedded subspaces. In *Proceedings of the 36th International Conference*

Table 1: Detailed results of high-dimensional optimization algorithms on synthetic functions with $D = 1000$. For methods based on subspace embedding techniques, we respectively evaluate two hyperparameter settings with $d = 30$ and $d = 50$.

Objective Function	Method	Convergence Value		Best Solution		Time (s)
		d=30	d=50	d=30	d=50	
Rosenbrock (D=1000)	SIRBO	413740.1875±120969.2969	459225.1875±96250.1875	259375.1562	348556.6562	54
	HesBO	633849.8750±200530.8281	819959.5000±137244.8906	341874.2500	681234.4375	125
	REMBO	6223221.1331±641472.7863	6183677.2826±809638.4103	5236903.2094	5851605.9492	766
	ALEBO	121149.1016±46185.1758	109807.8828±15043.3242	81243.7891	92437.4062	27665
	VAEBO	81353.2500±13184.3848	89912.2031±11958.9248	63902.0126	79479.9531	21107
	BxUS	33332.2138±2957.4265		23831.3776		1241
	MCTSVS	129055.41±59201.297		70404.7031		57
	TuRBO	79969.6875±14060.7158		56217.6797		156
	RDUCB	915660.5000±222155.5938		556649.5465		16653
	SBO-SE	643878.8125±201391.4688		254839.0312		1334
	DCEM	527429.5391±83910.8251		390377.3541		7
	LMMAES	447685.8176±110445.6733		252948.5421		14
	Random Search	2509598.2157±520699.5662		1677970.2069		4
	DSEBO (ours)	37699.6758±780.7019		37010.5820		1046
Sphere (D=1000)	SIRBO	87.9686±15.1527	95.3077±7.7289	49.4295	85.0751	90
	HesBO	120.2445±13.1439	114.5023±28.0399	99.4486	70.6100	305
	REMBO	350.8991±98.4352	498.2905±18.5448	192.5025	477.4872	453
	ALEBO	31.3825±10.4225	37.4376±2.5337	19.4079	35.9000	26648
	VAEBO	37.4775±6.6590	37.6364±5.9088	25.1240	31.1741	15163
	BxUS	17.8187±2.7321		14.4317		785
	MCTSVS	36.5522±4.2780		29.4209		51
	TuRBO	29.7788±11.4849		14.512		137
	RDUCB	174.8453±12.5177		151.0879		13759
	SBO-SE	152.2971±9.8699		135.7395		506
	DCEM	122.1613±15.5641		100.8576		11
	LMMAES	99.1157±15.992		76.3501		12
	Random Search	46.3139±37.8249		40.1777		5
	DSEBO (ours)	3.9387±1.3304		2.4966		339
Levy (D=1000)	SIRBO	90.0148±16.1821	97.3569±11.5434	63.7397	77.0712	95
	HesBO	149.0245±36.4953	146.9163±18.3832	117.1409	108.5397	295
	REMBO	1046.6322±67.1422	1206.4437±76.6944	927.1597	1087.9023	756
	ALEBO	35.7776±1.6821	35.0879±9.3454	34.5007	24.4679	25213
	VAEBO	18.6543±6.8300	19.4487±8.2072	11.1254	8.6804	15323
	BxUS	6.7886±0.0155		6.7753		630
	MCTSVS	32.6922±8.1272		18.0643		118
	TuRBO	6.4294±0.1092		6.2491		185
	RDUCB	198.5230±27.5693		146.3279		13496
	SBO-SE	184.6281±29.2152		110.4856		753
	DCEM	134.6127±14.4924		110.9705		9
	LMMAES	98.4986±9.0345		81.0548		11
	Random Search	113.8001±18.2524		88.0861		5
	DSEBO (ours)	2.2816±0.7567		1.5805		268
Griewank (D=1000)	SIRBO	61.7318±0.9419	62.0860±0.9048	60.1888	60.1376	92
	HesBO	64.8294±8.3552	73.6908±5.3223	44.9517	63.0216	315
	REMBO	11.4265±0.8796	220.8467±2.4297	11.0122	215.4218	639
	ALEBO	18.1015±1.3217	18.4540±1.2353	16.8223	17.3414	23792
	VAEBO	14.6719±0.4778	13.9477±0.6756	14.1418	12.7958	14685
	BxUS	11.2585±0.0474		11.1369		1360
	MCTSVS	62.4722±1.7076		60.0604		64
	TuRBO	63.1086±1.5921		61.1658		259
	RDUCB	164.7312±56.6559		74.8729		15583
	SBO-SE	88.7530±1.0904		86.7066		843
	DCEM	80.5171±1.4105		77.9347		10
	LMMAES	83.7417±2.1408		79.8172		12
	Random Search	41.1532±0.4551		40.3172		4
	DSEBO (ours)	11.2488±0.1092		10.9793		276
Dixon (D=1000)	SIRBO	582309.6250±127018.3516	512858.5000±120621.3672	399245.0938	282305.4062	116
	HesBO	512598.8125±150171.8906	662775.3750±159014.9062	314195.2812	361927.0312	148
	REMBO	10028901.2305±751474.3750	11296827.1480±432756.4375	8892939.2391	10683308.8750	531
	ALEBO	108987.9219±41572.4883	127629.6016±37048.0312	47649.1914	94961.3281	23283
	VAEBO	229556.6719±119026.5547	117513.2344±54179.7344	111616.2344	68923.3125	11638
	BxUS	44058.0664±15399.4717		26695.2305		953
	MCTSVS	126604.6875±52375.7891		56397.1484		58
	TuRBO	738974.3125±436899.8750		111616.2344		218
	RDUCB	1566930.6250±470125.1562		825860.3044		16561
	SBO-SE	1208742.0000±248711.9062		864583.7500		657
	DCEM	900571.3348±214454.3493		495204.757		11
	LMMAES	682351.8439±98667.3857		515423.2172		13
	Random Search	270112.6082±53654.5844		184623.3812		11
	DSEBO (ours)	39076.9609±10246.4609		25189.1094		597
Michalewicz (D=1000)	SIRBO	-8.4669±0.5782	-8.7477±0.7193	-9.3439	-9.8313	144
	HesBO	-8.3571±0.9363	-7.6276±0.9776	-9.9054	-9.1226	136
	REMBO	-3.3269±0.4194	-3.0310±0.4671	-3.9171	-3.7214	784
	ALEBO	-8.8393±0.5495	-8.4900±0.9351	-9.7144	-10.0362	26832
	VAEBO	-8.4725±1.0752	-7.7547±1.2024	-10.5696	-10.0460	11251
	BxUS	-10.1528±0.3488		-10.4914		1072
	MCTSVS	-7.9805±0.6546		-8.9399		34
	TuRBO	-8.4346±1.2460		-10.4085		246
	RDUCB	-6.4992±0.5495		-7.8331		14549
	SBO-SE	-7.7120±0.6833		-8.6905		542
	DCEM	-7.6842±0.8374		-9.2082		8
	LMMAES	-7.3345±0.4382		-7.9377		14
	Random Search	-7.7538±0.5373		-9.0669		16
	DSEBO (ours)	-10.6887±1.1657		-12.9019		465

Table 2: Detailed results of high-dimensional optimization algorithms on real-world tasks. For methods based on subspace embedding techniques, we respectively evaluate two hyperparameter settings with $d = 50$ and $d = 80$.

Objective Function	Method	Convergence Value		Best Solution		Time (s)
		d=50	d=80	d=50	d=80	
MSLR	SIRBO	-8.6166±0.0063	-8.5711±0.0094	-8.6594	-8.6050	71
	HesBO	-8.7637±0.0161	-8.6621±0.0272	-8.8077	-8.8044	219
	REMBO	-8.0890±0.0195	-8.8185±0.0145	-8.8886	-8.8850	252
	ALEBO	-8.8336±0.0112	-8.8620±0.0197	-8.8814	-8.8962	24016
	VAEBO	-7.7777±0.1623	-7.7285±0.1480	-8.5001	-8.2889	9766
	BAXUS	-8.8998±0.0964		-9.0185		554
	MCTSVS	-8.8755±0.1140		-9.0446		49
	TuRBO	-8.9199±0.1023		-9.0642		212
	SAASBO	-8.8604±0.0472		-8.9474		14326
	RDUCB	-8.8035±0.0944		-8.9027		3776
	DuMBO	-8.7808±0.1257		-8.9408		6873
	SBO-SE	-8.7492±0.1289		-8.9026		307
	DCEM	-8.7281±0.0806		-8.8316		5
	LMMAES	-8.6829±0.1933		-9.2051		8
	Random Search	-4.2603±1.4577		-7.5904		2
	DSEBO (ours)	-8.9396±0.0914		-9.1933		245
Lasso-Hard	SIRBO	39.8469±2.7875	40.0981±1.9785	37.9968	37.7477	145
	HesBO	41.2607±11.5219	47.6895±8.0687	24.9073	37.6182	317
	REMBO	31.1405±21.1006	41.2012±7.1255	6.7579	34.3955	740
	ALEBO	11.2730±5.7479	17.7599±3.1188	8.0080	12.9113	21846
	VAEBO	23.4285±9.4164	14.2004±3.6496	10.9416	9.8085	13657
	BAXUS	6.1045±0.5563		5.1150		642
	MCTSVS	9.4854±2.7239		7.6786		133
	TuRBO	11.5030±7.3585		4.2641		276
	RDUCB	11.6843±6.2225		4.4439		14924
	DuMBO	20.0604±4.6829		12.6193		14700
	SBO-SE	49.3089±4.2530		42.7787		584
	DCEM	43.6111±6.5621		32.1892		96
	LMMAES	5.5371±3.2662		1.8479		112
	Random Search	45.2683±1.5762		42.5421		29
	DSEBO (ours)	3.8613±0.4896		3.4248		483
LIMO	SIRBO	-5.2516±0.8961	-5.1352±0.4046	-6.8275	-5.8603	126
	HesBO	-5.2117±1.4180	-4.7430±0.8296	-7.7866	-6.2939	229
	REMBO	-3.526±0.9881	-3.6328±0.3971	-5.0612	-4.3895	164
	ALEBO	-7.5033±1.5741	-5.9202±0.7087	-10.4432	-6.8159	25496
	VAEBO	-8.0703±1.3382	-8.3918±2.0633	-9.7784	-11.4423	14205
	BAXUS	-6.8716±0.9177		-8.065		1562
	MCTSVS	-5.3277±0.7252		-6.5733		104
	TuRBO	-4.2479±1.3035		-6.9445		291
	RDUCB	-4.0750±0.6278		-5.4911		16256
	DuMBO	-5.9619±0.8675		-7.0755		21299
	SBO-SE	-4.6146±0.6944		-5.8465		544
	DCEM	-4.3841±0.4846		-5.3127		114
	LMMAES	-6.5043±1.1578		-7.805		126
	Random Search	-1.8629±0.6874		-3.0679		12
	DSEBO (ours)	-10.6613±2.4279		-14.2513		294

Table 3: Detailed results of high-dimensional optimization algorithms on synthetic functions with $D = 10000$. For methods based on subspace embedding techniques, we respectively evaluate two hyperparameter settings with $d = 30$ and $d = 50$.

Objective Function	Method	Convergence Value		Best Solution		Time (s)
		d=30	d=50	d=30	d=50	
Rosenbrock (D=10000)	SIRBO	654297.3750±251160.9062	620616.1250±226600.6094	391632.6250	429690.8438	7255
	HesBO	740927.6875±95669.8203	823125.9375±180349.7031	621726.5464	558716.9375	159
	REMBO	4762362.9526±1058705.0084	6043180.2348±539507.1027	3476606.4671	5216455.1280	795
	BAXUS	34236.7869±2550.1297		29355.3671		3995
	MCTSVS	57851.9000±6859.1123		50801.8242		272
	TuRBO	1788346.6±1056224.8		404563.2812		389
	DCEM	412348.2053±73896.1685		272307.0427		31
	LMMAES	572015.3508±114679.9362		357881.9302		54
	Random Search	3022829.2621±516710.6428		2262149.8444		18
	DSEBO (ours)	36658.1680±4038.2712		31940.3965		1076
Sphere (D=10000)	SIRBO	114.9471±17.8656	106.1076±12.5741	92.1035	85.4588	7528
	HesBO	143.2573±26.0684	135.1932±21.8210	113.4040	103.8779	230
	REMBO	282.7653±54.2532	497.3164±34.3274	198.8264	446.1657	581
	BAXUS	20.0685±2.4054		15.9891		3182
	MCTSVS	42.4199±11.5857		24.8201		383
	TuRBO	174.1312±39.0554		118.9279		264
	DCEM	124.469±14.4034		100.0146		30
	LMMAES	123.6035±14.0609		97.1999		50
	Random Search	68.6981±8.4385		53.9599		14
	DSEBO (ours)	6.4338±2.1632		4.6624		343
Levy (D=10000)	SIRBO	151.7623±8.6825	134.7558±19.8493	139.4174	112.8558	7632
	HesBO	176.8307±52.7794	147.1899±18.5760	119.6019	131.4840	155
	REMBO	1012.7800±231.5087	1266.2434±85.5595	606.2893	1168.6647	562
	BAXUS	57.4391±0.0222		57.4031		3038
	MCTSVS	64.1901±5.2330		58.3096		451
	TuRBO	72.8131±0.4172		72.2262		356
	DCEM	164.0851±22.484		126.7344		41
	LMMAES	140.5668±13.6518		120.7209		65
	Random Search	131.8982±18.8193		102.7094		17
	DSEBO (ours)	2.7522±1.2717		1.8678		290
Griewank (D=10000)	SIRBO	655.2772±17.9869	652.7079±12.9206	636.1684	643.0330	8631
	HesBO	571.2280±34.4433	711.1620±39.2822	522.5161	672.2633	157
	REMBO	101.1997±0.0365	2173.1267±9.0637	101.1423	2161.6895	537
	BAXUS	101.9942±0.5218		101.4823		3092
	MCTSVS	835.3915±6.4354		825.446		267
	TuRBO	882.3385±3.0383		878.386		382
	DCEM	830.4092±3.2544		825.6568		28
	LMMAES	840.1342±2.4772		835.4587		57
	Random Search	422.5575±1.2538		420.7823		16
	DSEBO (ours)	101.4063±0.0356		101.1153		527
Dixon (D=10000)	SIRBO	764620.5000±178713.5156	711741.6875±220006.9375	504689.5000	393210.5938	6855
	HesBO	527328.7500±293799.1562	593128.7500±238018.3594	195464.7500	247606.3281	195
	REMBO	8629979.1862±1030159.7500	9811402.4797±1358266.3750	7375704.5000	7227984.3564	658
	BAXUS	40101.4234±7970.3999		34711.2812		3463
	MCTSVS	288137.3750±118116.6797		138630.5781		250
	TuRBO	1712952.6250±153219.5156		1534449.3750		294
	DCEM	906037.6463±127875.9288		757046.0837		33
	LMMAES	967289.2009±145768.8284		658107.9851		55
	Random Search	277199.8422±75239.0032		107642.0207		19
	DSEBO (ours)	38253.0625±8405.7793		23499.7630		366
Michalewicz (D=10000)	SIRBO	-5.2707±0.8001	-5.4101±0.5353	-6.2494	-6.1494	5983
	HesBO	-5.3715±0.6174	-5.4323±0.6975	-6.4223	-6.5907	803
	REMBO	0.7975±0.2896	1.0713±0.1826	0.3354	0.8087	579
	BAXUS	-10.3466±1.5608		-11.8395		3992
	MCTSVS	-10.7175±1.1713		-11.5057		342
	TuRBO	-4.2744±0.8549		-5.4950		276
	DCEM	-5.2218±0.6693		-6.3025		37
	LMMAES	-4.6498±0.4431		-5.4561		62
	Random Search	-3.2545±0.4323		-3.9067		14
	DSEBO (ours)	-11.7390±0.6309		-12.7581		391

Table 4: Detailed results of MAB strategies on synthetic functions with $D = 1000$.

Objective Function	Method	Convergence Value	Best Solution	Time (s)
Rosenbrock (D=1000)	Extreme	2160763.7500 $\pm$ 1467483.8750	571516.5625	359
	Random	3700211.5000 $\pm$ 1493292.5000	1157037.0000	371
	ϵ -greedy	2273714.5000 $\pm$ 1358569.0000	52278.8203	1193
	C-UCB	2211636.7500 $\pm$ 1624777.0000	779282.4375	395
	Softmax	2211636.7500 $\pm$ 1624777.0000	779282.4375	398
	TS	1856110.0000 $\pm$ 834666.1875	383319.5625	474
	UCB-E	2248272.2500 $\pm$ 1751758.2500	779282.4375	463
	S-Halving	2328358.7500 $\pm$ 1153945.5000	463237.6562	401
	Expectation	2248154.2500 $\pm$ 1751778.3750	779282.4375	555
	DSEBO (ours)	37699.6758$\pm$780.7019	37010.5820	1046
Sphere (D=1000)	Extreme	129.6299 $\pm$ 181.8460	13.1881	1223
	Random	150.9959 $\pm$ 121.5409	25.4365	146
	ϵ -greedy	19.5962 $\pm$ 3.7292	13.1885	468
	C-UCB	22.8336 $\pm$ 7.7990	13.1881	1603
	Softmax	63.3230 $\pm$ 127.7900	13.1881	1348
	TS	20.9170 $\pm$ 4.5001	13.1881	1791
	UCB-E	22.8336 $\pm$ 7.7990	13.1881	1555
	S-Halving	80.7406 $\pm$ 113.8652	13.1898	939
	Expectation	20.5691 $\pm$ 4.5567	13.1881	1613
	DSEBO (ours)	3.9387$\pm$1.3304	2.4966	399
Levy (D=1000)	Extreme	363.8994 $\pm$ 380.7985	4.2985	1013
	Random	750.0236 $\pm$ 276.8358	40.4695	134
	ϵ -greedy	217.8715 $\pm$ 167.7238	2.9233	411
	C-UCB	233.7007 $\pm$ 167.9456	4.2985	1289
	Softmax	272.7462 $\pm$ 251.0918	4.2985	1393
	TS	267.9188 $\pm$ 218.1308	4.2985	1288
	UCB-E	217.7698 $\pm$ 150.8256	4.2985	1268
	S-Halving	360.4880 $\pm$ 277.0266	18.9699	976
	Expectation	222.1218 $\pm$ 201.5956	4.2985	1227
	DSEBO (ours)	2.2816$\pm$0.7567	1.5805	268
Griewank (D=1000)	Extreme	11.4613 $\pm$ 0.3793	11.2658	1285
	Random	11.8220 $\pm$ 1.5041	11.2802	95
	ϵ -greedy	11.5547 $\pm$ 0.6718	11.2663	365
	C-UCB	11.4613 $\pm$ 0.3793	11.2658	1285
	Softmax	11.4613 $\pm$ 0.3793	11.2658	1289
	TS	11.4613 $\pm$ 0.3793	11.2658	1287
	UCB-E	11.4613 $\pm$ 0.3793	11.2658	1290
	S-Halving	11.6494 $\pm$ 1.0168	11.2120	1047
	Expectation	11.4613 $\pm$ 0.3793	11.2658	1301
	DSEBO (ours)	11.2488$\pm$0.1092	10.9793	276
Dixon (D=1000)	Extreme	6540823.5000 $\pm$ 2834092.7500	1504931.8750	455
	Random	6720030.5000 $\pm$ 1582640.6250	4071802.5000	592
	ϵ -greedy	5106037.5000 $\pm$ 1698677.7500	2753771.2500	362
	C-UCB	4299205.5000 $\pm$ 2713686.0000	51601.1992	465
	Softmax	4299205.5000 $\pm$ 2713686.0000	51601.1992	604
	TS	3428899.5000 $\pm$ 2103435.0000	33287.6367	621
	UCB-E	4299205.5000 $\pm$ 2713686.0000	51601.1992	716
	S-Halving	6044710.5000 $\pm$ 3704331.5000	2606495.7500	1124
	Expectation	4299205.5000 $\pm$ 2713686.0000	51601.1992	603
	DSEBO (ours)	39076.9609$\pm$10246.4609	25189.1094	597
Michalewicz (D=1000)	Extreme	-4.6286 $\pm$ 2.2142	-8.8314	503
	Random	-4.2663 $\pm$ 0.8069	-5.8979	92
	ϵ -greedy	-5.7135 $\pm$ 1.5460	-7.9152	271
	C-UCB	-6.0991 $\pm$ 0.9943	-7.8918	366
	Softmax	-6.0994 $\pm$ 0.9661	-7.3633	570
	TS	-6.6686 $\pm$ 1.2886	-8.8314	284
	UCB-E	-6.2666 $\pm$ 1.0264	-7.5300	204
	S-Halving	-4.8470 $\pm$ 1.2241	-6.6878	429
	Expectation	-6.5533 $\pm$ 1.3236	-8.8314	344
	DSEBO (ours)	-10.6887$\pm$1.1657	-12.9019	465

Table 5: Detailed results of MAB strategies on real-world tasks.

Objective Function	Method	Convergence Value	Best Solution	Time (s)
MSLR	Extreme	-8.8175 $\pm$ 0.0982	-8.9992	548
	Random	-8.8245 $\pm$ 0.1437	-9.0390	129
	ϵ -greedy	-8.7972 $\pm$ 0.0639	-8.9243	346
	C-UCB	-8.8482 $\pm$ 0.0580	-8.9549	494
	Softmax	-8.8766 $\pm$ 0.0657	-9.0177	663
	TS	-8.8586 $\pm$ 0.0928	-9.0130	744
	UCB-E	-8.9054 $\pm$ 0.1266	-9.1812	332
	S-Halving	-8.8690 $\pm$ 0.1694	-9.1396	557
	Expectation	-8.8185 $\pm$ 0.1308	-9.1285	627
	DSEBO (ours)	-8.9396$\pm$0.0914	-9.1933	245
Lasso-Hard	Extreme	52.0638 $\pm$ 19.5580	7.6819	402
	Random	52.3876 $\pm$ 15.1716	15.5185	342
	ϵ -greedy	44.1027 $\pm$ 17.8642	7.4094	438
	C-UCB	43.4808 $\pm$ 23.4663	7.4081	819
	Softmax	45.7185 $\pm$ 25.1736	7.4081	821
	TS	39.6904 $\pm$ 24.0790	6.7548	783
	UCB-E	46.0051 $\pm$ 22.9438	7.6461	382
	S-Halving	53.6740 $\pm$ 14.7570	27.8280	405
	Expectation	42.9734 $\pm$ 21.6694	7.4081	795
	DSEBO (ours)	3.8613$\pm$0.4896	3.4248	483
LIMO	Extreme	-3.9495 $\pm$ 1.0746	-5.7674	406
	Random	-3.7930 $\pm$ 1.0581	-5.2494	161
	ϵ -greedy	-4.0860 $\pm$ 0.8866	-6.0472	276
	C-UCB	-5.1490 $\pm$ 1.6835	-8.1142	151
	Softmax	-4.4801 $\pm$ 1.1440	-6.0536	433
	TS	-4.4772 $\pm$ 0.7694	-5.7674	386
	UCB-E	-4.8374 $\pm$ 1.2162	-6.6388	789
	S-Halving	-3.9992 $\pm$ 1.1429	-6.6801	193
	Expectation	-4.0141 $\pm$ 1.5697	-6.2238	376
	DSEBO (ours)	-10.6613$\pm$2.4279	-14.2513	294

- on *Machine Learning*, volume 97, pages 4752–4761, Long Beach, CA, 2019.
- Leonard Papenmeier, Luigi Nardi, and Matthias Poloczek. Increasing the scope as you learn: Adaptive Bayesian optimization in nested subspaces. In *Advances in Neural Information Processing Systems 35*, pages 11586–11601, New Orleans, LA, 2022.
- Hong Qian, Yi-Qi Hu, and Yang Yu. Derivative-free optimization of high-dimensional non-convex functions by sequential random embeddings. In *Proceedings of the 25th International Joint Conference on Artificial Intelligence*, pages 1946–1952, New York, NY, 2016.
- Tao Qin, Tie-Yan Liu, Jun Xu, and Hang Li. LETOR: A benchmark collection for research on learning to rank for information retrieval. *Information Retrieval*, 13(4):346–374, 2010.
- Kenan Sehic, Alexandre Gramfort, Joseph Salmon, and Luigi Nardi. Lassobench: A high-dimensional hyperparameter optimization benchmark suite for lasso. In *Proceedings of the 2nd International Conference on Automated Machine Learning*, volume 188, pages 2/1–24, Baltimore, MD, 2022.
- Bobak Shahriari, Kevin Swersky, Ziyu Wang, Ryan P. Adams, and Nando de Freitas. Taking the human out of the loop: A review of Bayesian optimization. *Proceedings of the IEEE*, 104(1):148–175, 2016.
- Lei Song, Ke Xue, Xiaobin Huang, and Chao Qian. Monte Carlo tree search based variable selection for high dimensional Bayesian optimization. In *Advances in Neural Information Processing Systems 35*, pages 28488–28501, New Orleans, LA, 2022.
- Niranjan Srinivas, Andreas Krause, Sham M. Kakade, and Matthias W. Seeger. Gaussian process optimization in the bandit setting: No regret and experimental design. In *Proceedings of the 27th International Conference on Machine Learning*, pages 1015–1022, Haifa, Israel, 2010.
- Tianxiang Sun, Zhengfu He, Hong Qian, Yunhua Zhou, Xuanjing Huang, and Xipeng Qiu. Bbtv2: Towards a gradient-free future with large language models. In *Proceedings of the 27th Conference on Empirical Methods in Natural Language Processing*, pages 3916–3930, Abu Dhabi, United Arab Emirates, 2022.
- Richard S Sutton and Andrew G Barto. *Reinforcement Learning: An Introduction*. MIT press, Cambridge, 2018.
- Ziyu Wang, Frank Hutter, Masrour Zoghi, David Matheson, and Nando de Freitas. Bayesian optimization in a billion dimensions via random embeddings. *Journal of Artificial Intelligence Research*, 55:361–387, 2016.
- Zhitong Xu, Haitao Wang, Jeff M Phillips, and Shandian Zhe. Standard Gaussian process is all you need for high-dimensional bayesian optimization. In *Proceedings of the 13th International Conference on Learning Representations*, 2025.
- Miao Zhang, Huiqi Li, and Steven W. Su. High dimensional Bayesian optimization via supervised dimension reduction. In *Proceedings of the 28th International Joint Conference on Artificial Intelligence*, pages 4292–4298, Macao, China, 2019.
- Yangwenhui Zhang, Hong Qian, Xiang Shu, and Aimin Zhou. High-dimensional dueling optimization with preference embedding. In *Proceedings of the 37th AAAI Conference on Artificial Intelligence*, pages 11280–11288, Washington, DC, 2023.
- Juliusz Krzysztof Ziomek and Haitham Bou-Ammar. Are random decompositions all we need in high dimensional Bayesian optimisation? In *Proceedings of the 40th International Conference on Machine Learning*, volume 202, pages 43347–43368, Honolulu, HI, 2023.